

# Docker Security Workshop

David Lawrence

Riyaz Faizullabhoy

Viktor Stanchev

# Goals of this workshop

# Features

- See what features are available in Docker
- What do they do?
- How do you use them?

# Understanding

- Look at some tools
- See underlying implementation details
- Learn best practices

# Do!

APPROVED

**Don't!**

**CAUTION!**

Docker is additive to the security of your application ...

... even if you don't use any of the  
techniques we cover



# Docker aims to be Secure by Default

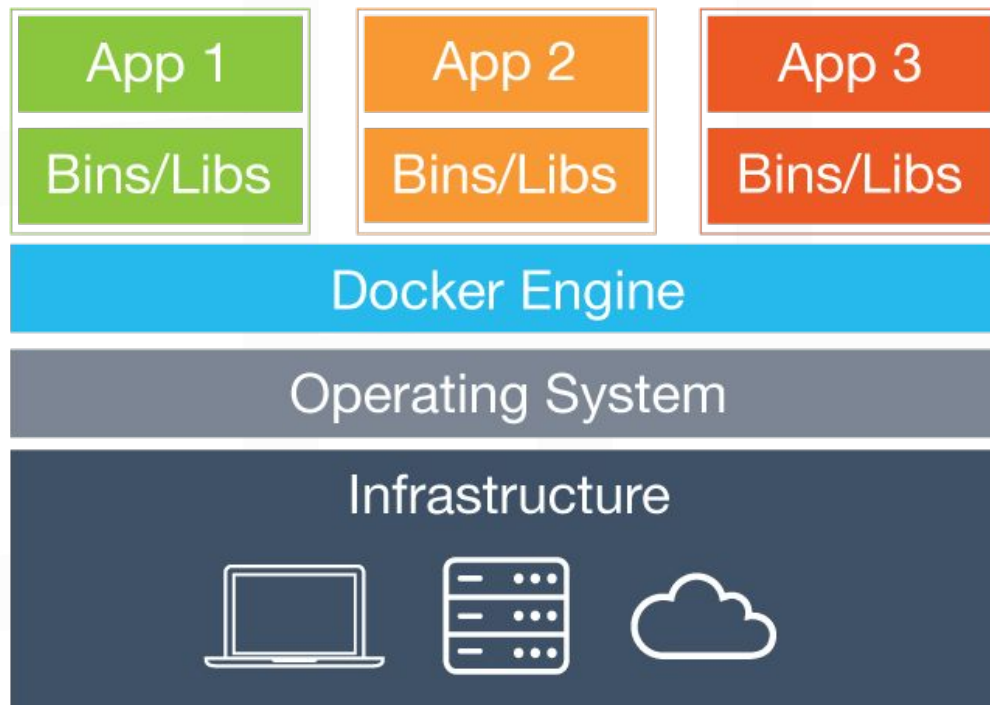
<https://docs.docker.com/engine/security/non-events/>

CVE-2013-1956, 1957, 1958, 1959, 1979, CVE-2014-4014, 5206, 5207, 7970, 7975, CVE-2015-2925, 8543, CVE-2016-3134, 3135, CVE-2014-0181, CVE-2015-3339, CVE-2014-4699, CVE-2014-9529, CVE-2015-3214, 4036, CVE-2016-0728, CVE-2016-2383

# Why?


# How do we think about containers?

# How we talk about Docker



# How Docker Actually Works

Host OS



Where can we see this?

top

# Anatomy of a Container

# Namespaces: “what containers can see”

```
ls -la /proc/<pid>/ns/
```



# What's namespaced?

Cgroup	CLONE_NEWCGROUP	Cgroup root directory
IPC	CLONE_NEWIPC	System V IPC, POSIX message queues
Network	CLONE_NEWNET	Network devices, stacks, ports, etc.
Mount	CLONE_NEWNS	Mount points
PID	CLONE_NEWPID	Process IDs
User	CLONE_NEWUSER	User and group IDs
UTS	CLONE_NEWUTS	Hostname and NIS domain name

# Demo using namespaces directly

Create a shell process with pid and fs namespaces.

```
$ sudo unshare -fp
```


```
$ sudo unshare -fp --mount-proc
```

# cgroups: “what containers can use”


Aka “Control Groups” - limit container resources!

- CPU
- Memory
- PIDs

# cgroups: “what containers can use”




# cgroups: “what containers can use”



cgroups: “what containers can use”

: ( ) { : | : & } ; :



# Hands-On Exercise

Set up your AWS instance - check your email!

```
chmod 400 <PATH_TO_FILE>/<name>.pem  
ssh -i <PATH_TO_FILE>/<name>.pem ubuntu@<Public DNS>
```

**Example:** `ssh -i riyaz.pem ubuntu@ec2-54-149...compute.amazonaws.com`

---

```
git clone https://github.com/riyazdf/dockercon-workshop.git  
- cgroups directory
```


**NOTE:** Ubuntu 15.10 does not support PID limits, but 16.04 does if you have it  
So *DO NOT* run the fork bomb unless you have another machine.

# Securing Client-Engine Communications

“My first 5 minutes...”



# Docker Client Server Architecture



# Exposing your engine to the internet

Edit config at `/lib/systemd/system/docker.service`

```
- ExecStart=/usr/bin/docker daemon -H fd://
```

```
+ ExecStart=/usr/bin/docker daemon -H fd:// -H tcp://0.0.0.0:2376
```

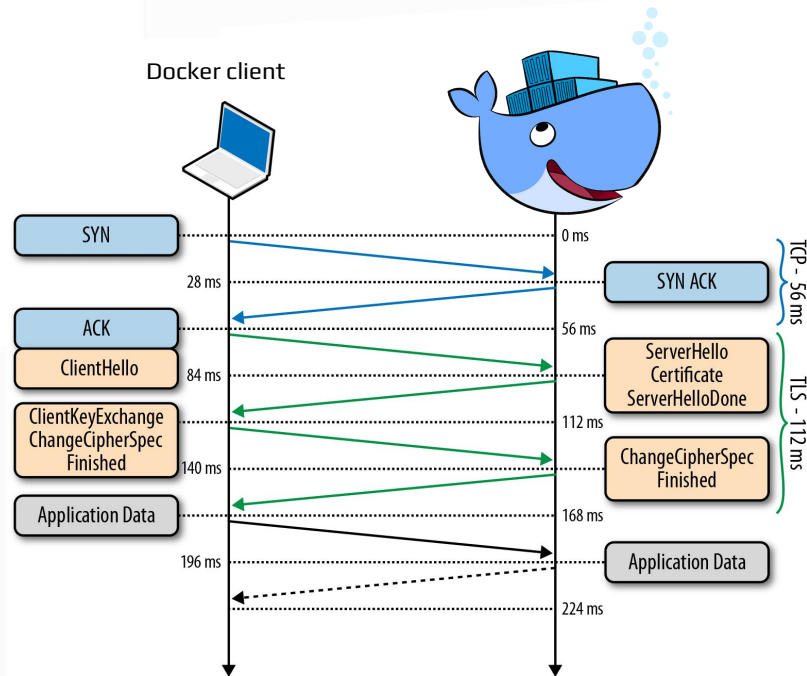
Restart Docker

```
$ sudo systemctl daemon-reload
```

```
$ sudo systemctl restart docker
```

# One Way TLS

- Same way we trust websites:
  - Server cert and key on engine
  - CA cert on client
  - client authenticates Docker engine



# Creating a CA

# use a strong passphrase!

```
$ openssl genrsa -aes256 -out ca-key.pem 4096
```

```
$ openssl req -new -x509 -days 365 -key ca-key.pem -sha256 -out ca.pem
```

# Creating the daemon cert and key

```
$ openssl genrsa -out server-key.pem 4096
```

```
$ openssl req -subj "/CN=$HOSTNAME" -sha256 -new -key server-key.pem \  
-out server.csr
```

```
$ echo subjectAltName = IP:10.10.10.20,IP:127.0.0.1 > extfile.cnf
```

```
$ openssl x509 -req -days 365 -sha256 -in server.csr -CA ca.pem \  
-CAkey ca-key.pem -CAcreateserial -out server-cert.pem \  
-extfile extfile.cnf
```

# Starting the daemon with the cert and key

```
$> tree /etc/docker
```

```
├── key.json  
├── server.pem  
└── server-key.pem
```

```
$> /usr/bin/docker daemon \
```

```
-H tcp://0.0.0.0:2376 -H unix:///var/run/docker.sock \ --storage-driver aufs \  
--tlsverify \  
--tlscert /etc/docker/server.pem \  
--tlskey /etc/docker/server-key.pem
```

# Trusting the daemon's cert on the client

```
$> tree ~/.docker
```

```
├── config.json  
└── ca.pem
```

```
$> export DOCKER_CERT_PATH=~/.my_cert_directory
```

```
$> tree ~/.my_cert_directory
```

```
└── ca.pem
```

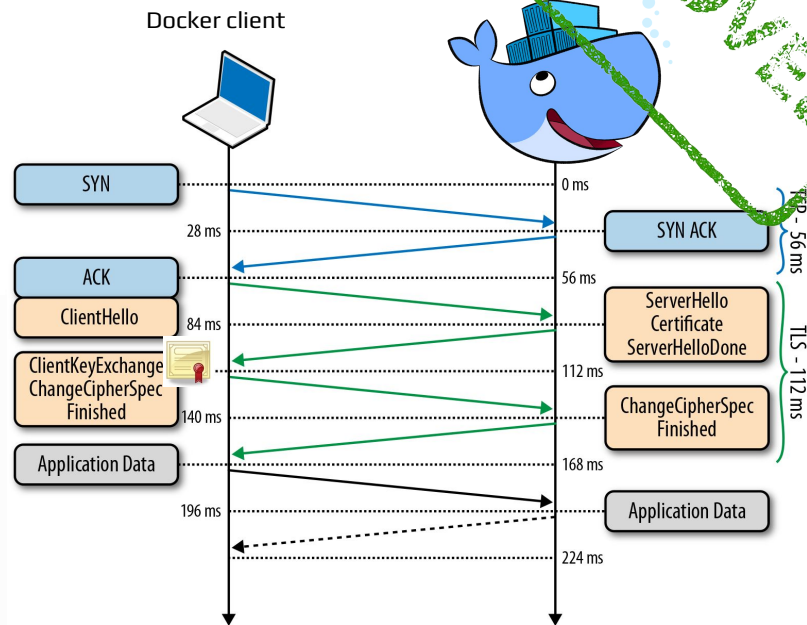
# Secure by default: docker-machine

docker-machine does this automatically to set up TLS for you by default!



# Best practice: Mutual TLS

- Client also presents certificate
  - Sends after verifying server cert
  - Mutual authentication
- Client CA on daemon (engine)



# Creating client cert and key

```
$ openssl genrsa -out key.pem 4096
```

```
$ echo extendedKeyUsage = clientAuth > extfile.cnf
```

```
$ openssl x509 -req -days 365 -sha256 -in client.csr -CA ca.pem \  
-CAkey ca-key.pem -CAcreateserial -out cert.pem -extfile extfile.cnf
```

NOTE: this `ca.pem` can (and should) be a different CA

# Trusting the client cert on the daemon

```
$> tree /etc/docker
```

```
├── key.json  
├── server.pem  
├── server-key.pem  
└── ca.pem
```

```
$> /usr/bin/docker daemon \
```

```
-H tcp://0.0.0.0:2376 -H unix:///var/run/docker.sock \ --storage-driver aufs \  
--tlsverify \  
--tlscert /etc/docker/server.pem \  
--tlskey /etc/docker/server-key.pem  
--tlscacert /etc/docker/ca.pem
```

# Using the client certs on the client

```
$> tree ~/.docker
```

```
|— config.json  
|— ca.pem  
|— cert.pem  
|— key.pem
```

```
$> export DOCKER_CERT_PATH=~/.my_cert_directory
```

```
$> tree ~/.my_cert_directory
```

```
|— ca.pem  
|— cert.pem  
|— key.pem
```

# Securing Engine-Registry Communications

/etc/docker/certs.d/	<-- Certificate directory
└─ localhost	<-- Hostname
└─ client.cert	<-- Client certificate
└─ client.key	<-- Client key
└─ localhost.crt	<-- Certificate authority that signed the registry certificate

# What's in an image: The Layered Filesystem

# What is a layered filesystem?

Combine multiple directories to look like a single filesystem

- Tombstoning/whiteout files to delete files from lower layers



# Supported Implementations

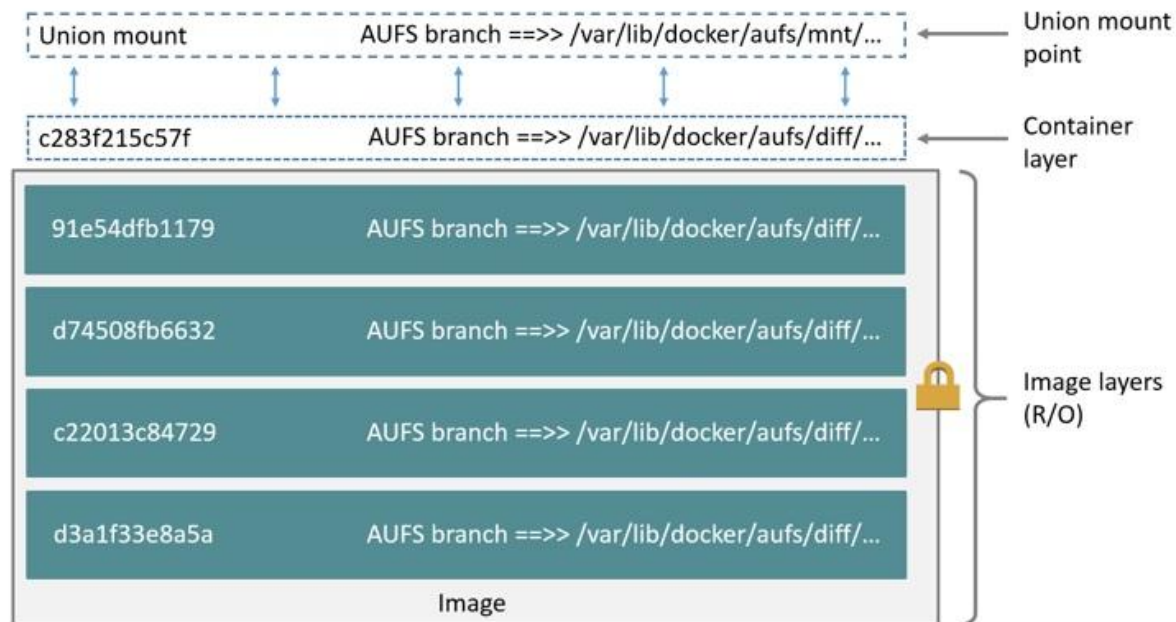
Aufs

Btrfs

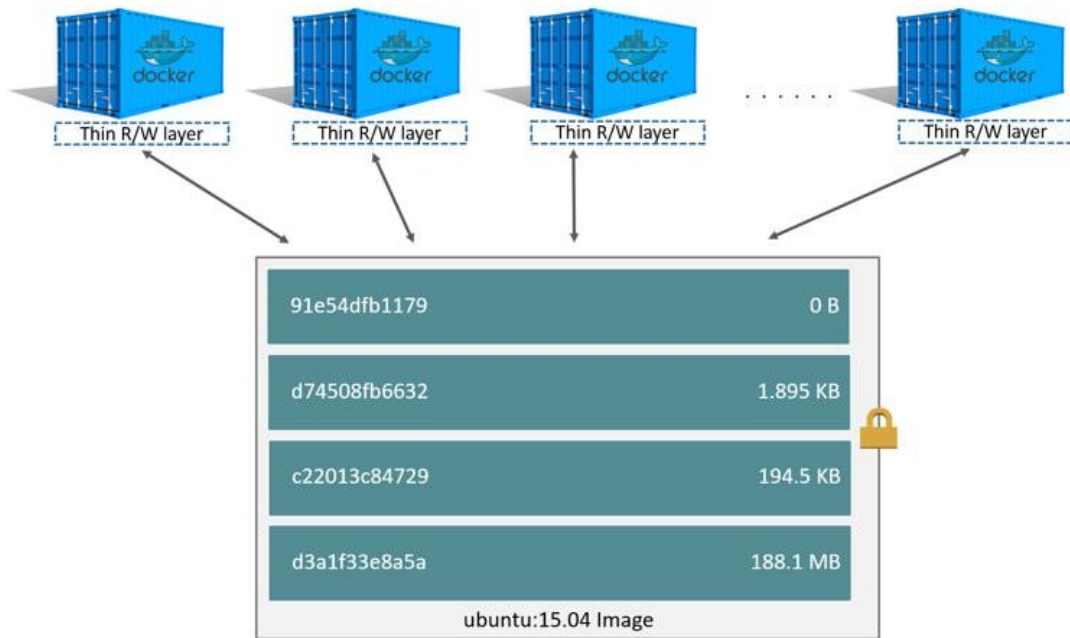
OverlayFS

Devicemapper

...



# Copy-on-write



# Best practice: “minimal” base images



alpine

- ~ 2 MB from hub (1 layer!)
- musl libc and busybox

ubuntu

- ~ 50 MB from hub

# Best practice: verify content



```
RUN apt-key adv \
    --keyserver hkp://keyserver.ubuntu.com:80 \
    --recv-keys BBEBDCB318AD50EC6865090613B00F1FD2C19886 \
    && echo deb http://repository.spotify.com stable non-free \
    | sudo tee /etc/apt/sources.list.d/spotify.list
```

# Best practice: read only containers



```
$ docker run it --rm --read-only alpine sh
```

Mounts the container's FS as read-only

# Best practice: read-only Volumes



```
-v /data:/data:ro
```

**Common mistake:**  
mount host location as writable



```
$ docker run it --rm -v /:/host alpine sh
```

## Best practice: minimal, read-only mounts



```
$ docker run it --rm -v /subdir/we/need:/dir:ro alpine sh
```



# Networks

# Isolate services

Control which services can talk to which other services

- Easier to audit

# Links (legacy)

Allow 2 specific containers to talk to each other.

- Brittle: does not survive container restarts


```
docker run -d --name db mysql:latest  
docker run -d --link db wordpress
```

# Network Namespace

```
docker network create my_app  
docker run -it --rm --net=my_app alpine sh
```

Links are dynamic, can be created to not yet created containers.

# Best practice: Use Multiple Networks



# Common Mistake: `--net=host`



Container can see ***ALL*** network traffic, including traffic on docker virtual networks

# Common Mistake: ports exposed on host



- Unnecessary
- Creates conflicts

# Best practice: Mutual TLS

*Implementation detail:* use mutual TLS between pairs of services that need to talk to each other.





# User Management

# Default runs as root

```
$ docker run -v /bin:/host/bin -it --rm alpine sh
/ $ whoami
root
/ $ id
uid=0(root) gid=0(root)
/ $ rm /host/bin/sh # WREAK HAVOC TIME! Please don't do this
```

# root in container == root outside container



0 (root)

1000 (user)

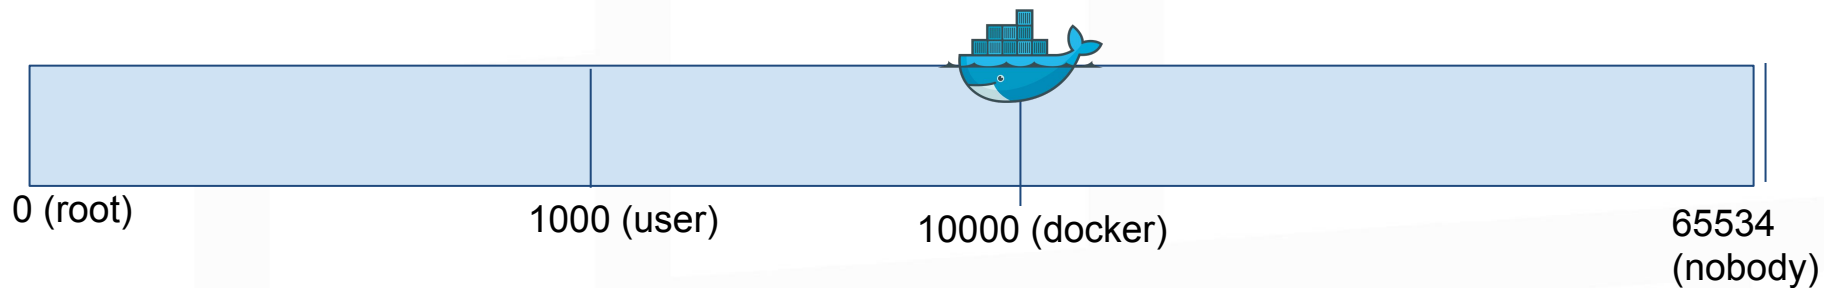
65534  
(nobody)

We don't want this to be the case!  
How can we change this?

# Step in the right direction: run as a user

```
# Use the --user flag with UID:GID argument
$ docker run -v /bin:/host/bin --user 10000:10000 -it --rm alpine sh
/ $ whoami
whoami: unknown uid 10000
/ $ id
uid=10000 gid=10000
/ $ rm /host/bin/sh
rm: can't remove 'sh': Permission denied
```

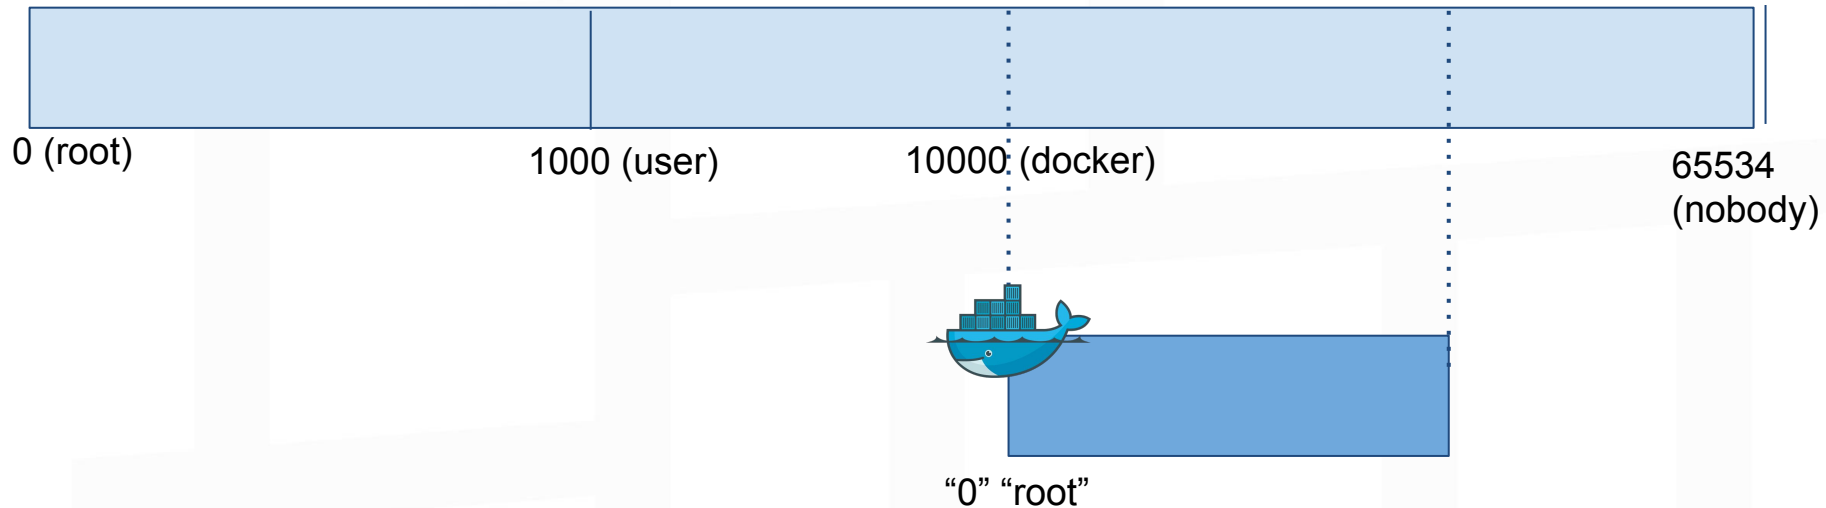
# But I still want “root” inside container



Perhaps we need to run a command that needs to look like it's root in the container, but we don't want to give it *true* root access to the underlying host

# Enable user namespaces

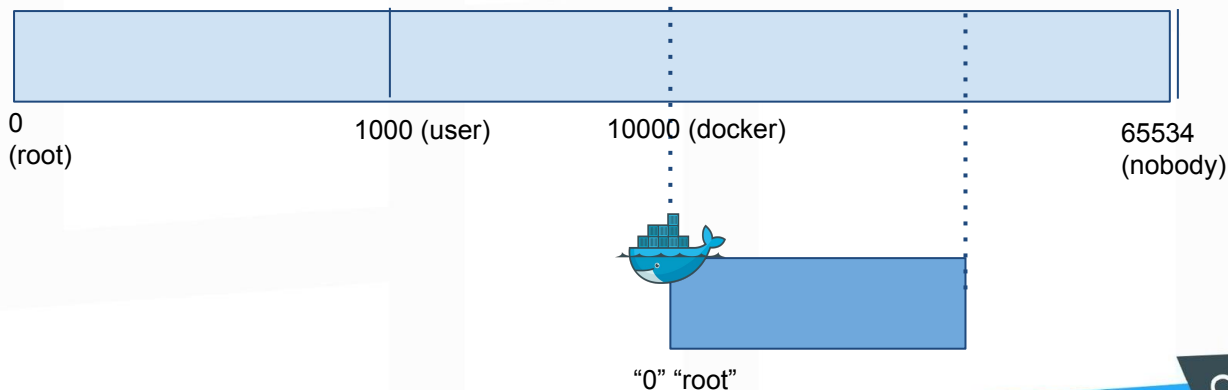
```
$ docker daemon --userns-remap [uid[:gid]]
```



# Enable user namespaces - common pitfalls

**\$ docker daemon --userns-remap [uid[:gid]]**

- Will need to re-pull images and re-create volumes due to container resource and image layer permissions
  - Leave this feature on in production; switching back and forth should only be done in development



# Hands-On Exercise & break

[github.com/riyazdf/dockercon-workshop](https://github.com/riyazdf/dockercon-workshop) - `usersns` directory



# Image Distribution

# Security Goals

## Image *Provenance* and *Trust*

- **Provenance:** who made this image?
  - Verify the *publisher* of the image
- **Trust:** have the contents of this image been tampered with?
  - Verify the *integrity* of the image

# Pulling by tag

```
$ docker pull alpine:latest
```

Name resolution takes place in registry to find content-address of image

```
$ docker pull alpine
```

Using default tag: latest

Notice that the tag defaults to latest if no tags are given!

# Pulling by digest

```
$ docker pull alpine@sha256:ea0d1389812...
```

No name resolution!

*Security best practice:* pulling by digest to enforce consistent and “immutable” pulls because of content-addressability

# Content Trust


```
$ export DOCKER_CONTENT_TRUST=1
```

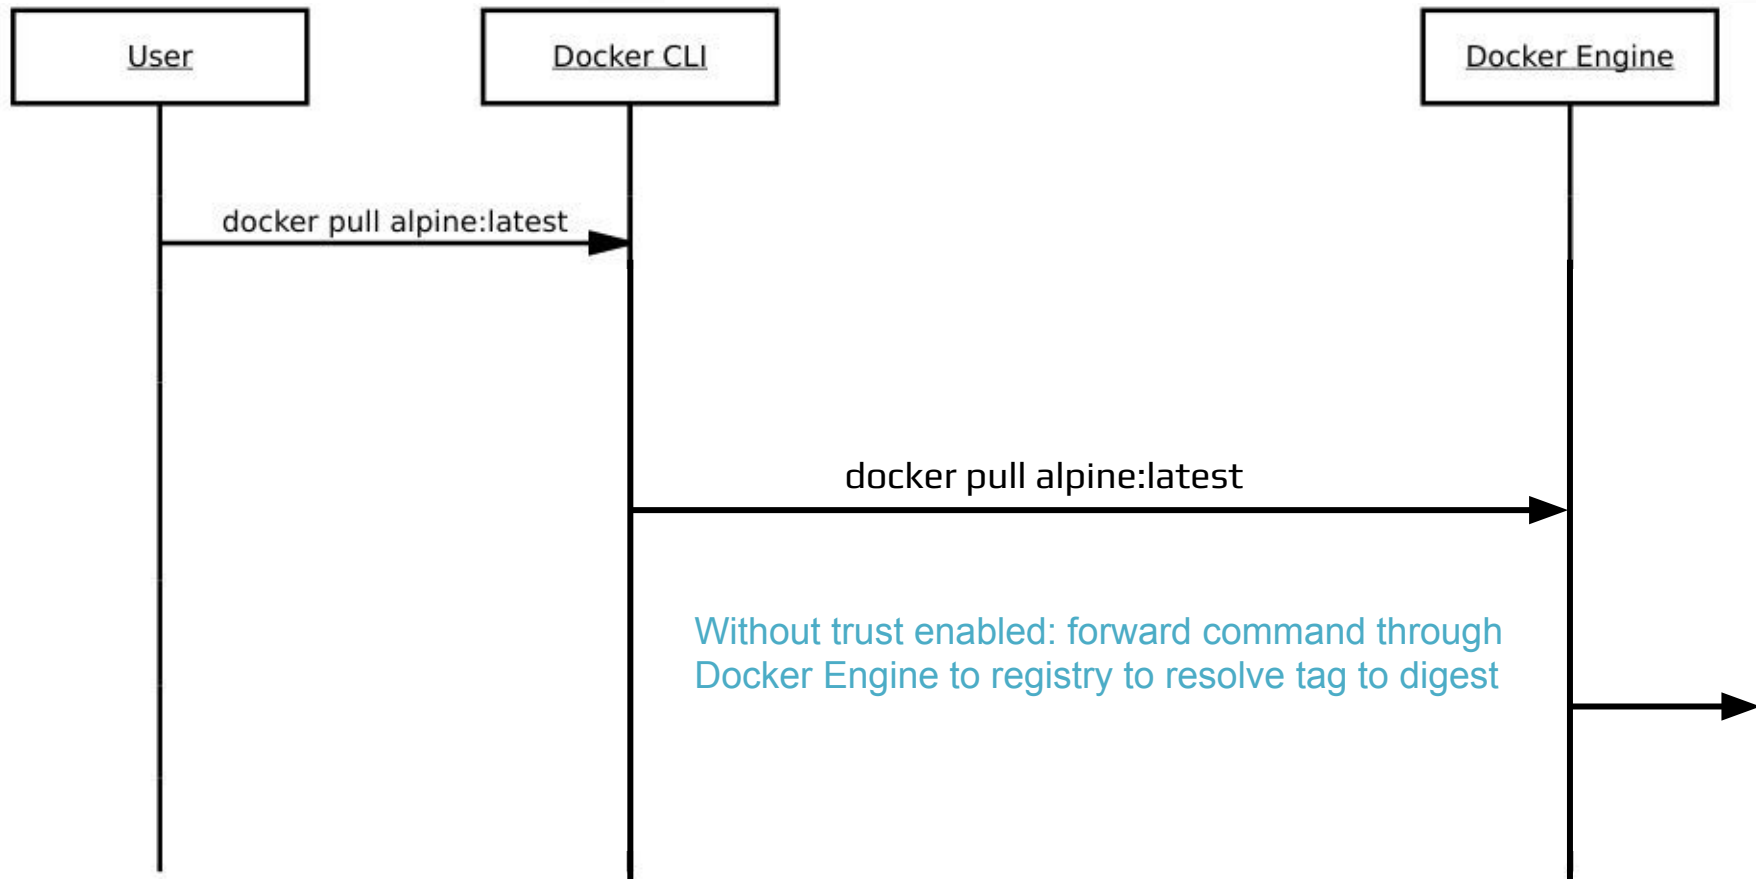
```
$ docker pull alpine:latest
```

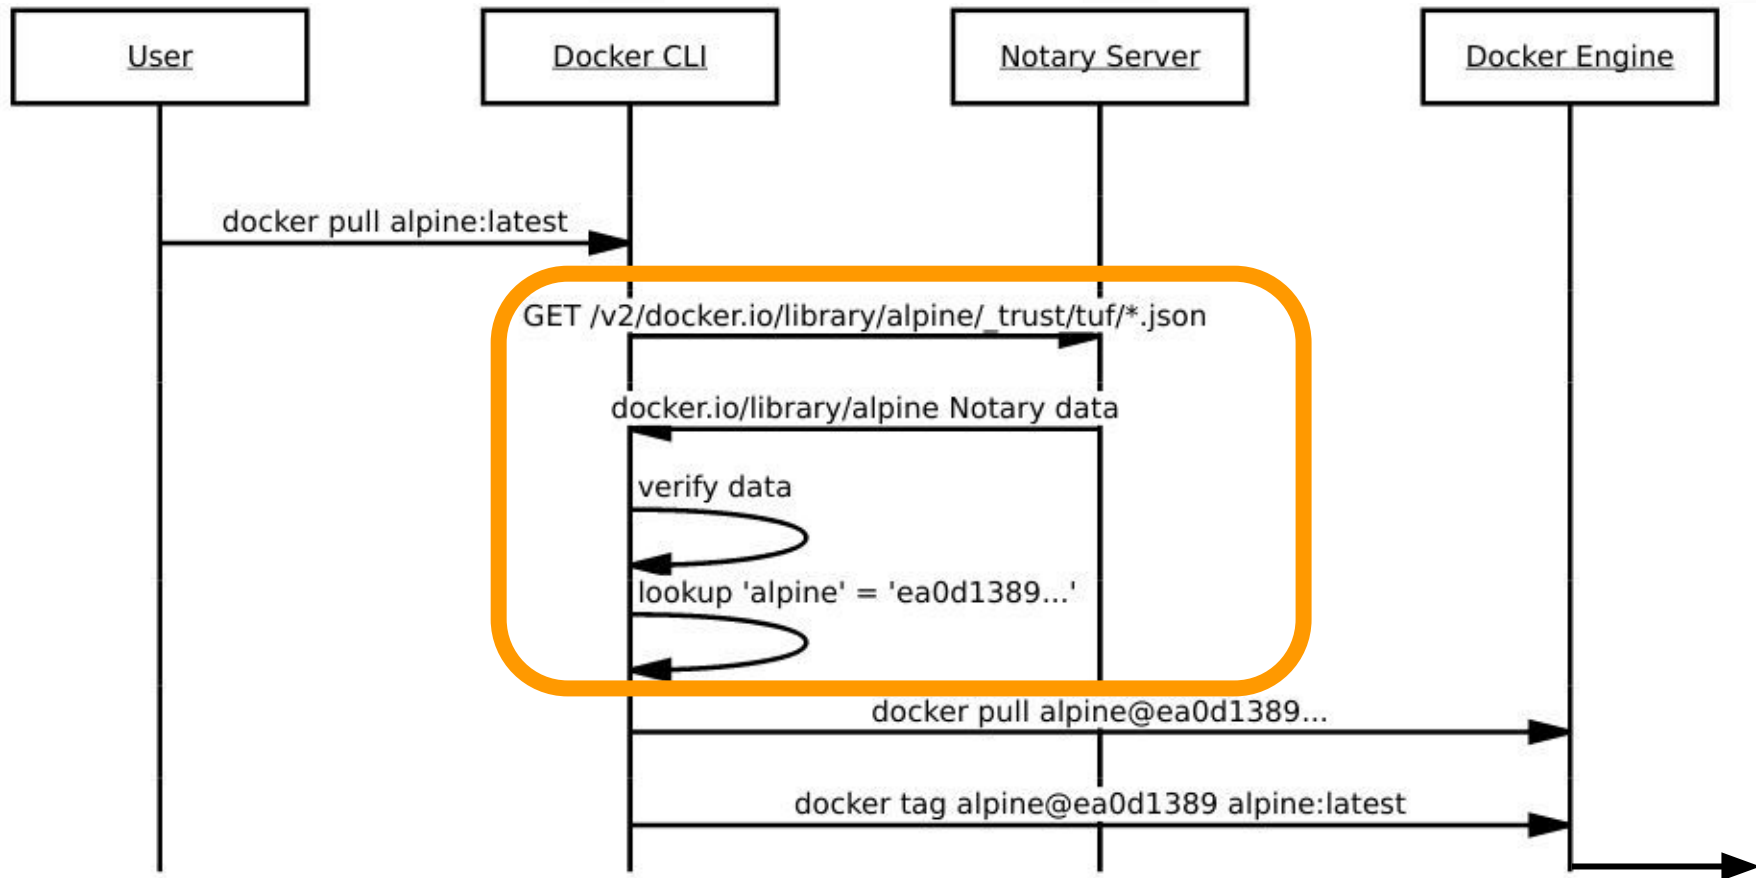
```
Pull (1 of 1): alpine:latest@sha256:ea0d1389...
```

Benefits of pull by digest with ease of pull by tag

APPROVED








# Content Trust (on push)

```
$ export DOCKER_CONTENT_TRUST=1  
$ docker tag alpine:latest <user>/alpine:trust  
$ docker push <user>/alpine:trust
```

Looks the same as a regular push by tag!





# Content Trust (it's more than gpg)

The push refers to a repository [<user>/alpine]

77f08abee8bf: Pushed

trust: digest: sha256:d5de850d728... size: 1355

Signing and pushing trust metadata

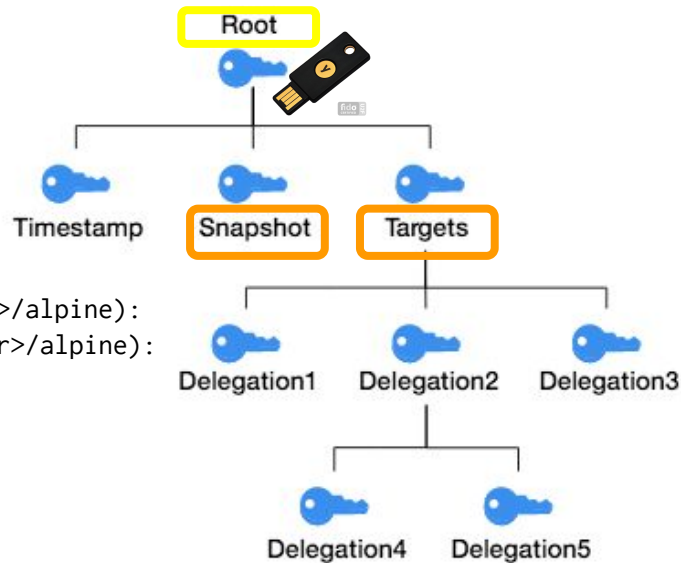
Enter passphrase for root key with ID e83f424:

Enter passphrase for new repository key with ID f903fc9 (docker.io/<user>/alpine):

Repeat passphrase for new repository key with ID f903fc9 (docker.io/<user>/alpine):

Finished initializing "docker.io/<user>/alpine"

Successfully signed "docker.io/<user>/alpine":trust



# Content Trust (it's more than gpg)

```
{
  "signed": {
    "expires": "2016-05-10T17:41:03.201245515Z",
    "snapshot": {
      "hashes": {
        "sha256": "qDEr1jJYkRjQkpN7RBYBDp15EJptU4vmFto707reX0M="
      },
      "length": 1545
    }
  },
  "version": 24
},
"signatures": [
  {
    "keyid": "b7c59624ccb68326737b34fc7ad4256d491cd50dbe64b958ab617a571607271d",
    "method": "ecdsa",
    "sig": "mLYqNVdMIAPXOLXJIj14AVqduP7bZGKH+7010mJOJ3z84s9xrBPzwdcp3SrFSstpaNZa5RuHpF++XjKJl1BfuA=="
  }
]
}
```



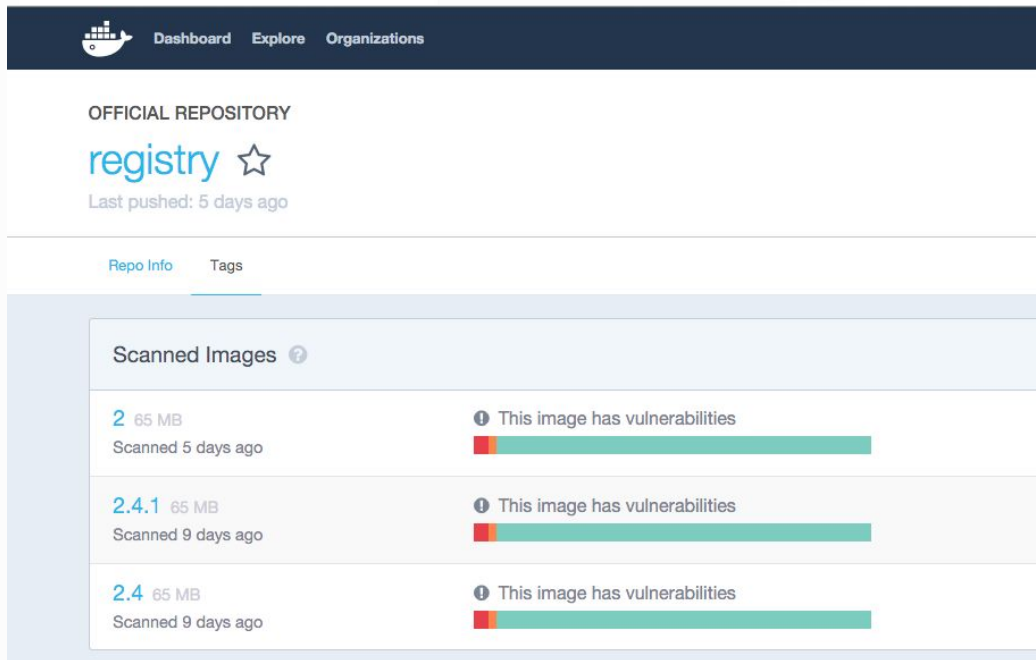
```
$ cat ~/.docker/trust/tuf/docker.io/alpine/metadata/timestamp.json | jq
```

# Docker Content Trust / Notary Threat Model

- **Key compromise?**
  - We can recover!
- **Replay attacks?**
  - Not with our freshness guarantees!
- **Untrusted registry?**
  - No problem! DCT/Notary do not root any trust in the underlying content store or transport
  - Use signed TUF metadata to retrieve trusted hashes of content
  - Don't even need to trust Notary server after first pull - local metadata pins trust, tagging keys are kept client-side for signing

# Docker Pull

Only pull trusted images  
Use official images when possible!




The screenshot shows the Docker Registry interface. At the top, there's a dark blue navigation bar with the Docker logo and links for Dashboard, Explore, and Organizations. Below this, the page is titled "OFFICIAL REPOSITORY" and "registry" with a star icon. It indicates "Last pushed: 5 days ago". The main content area has two tabs: "Repo Info" and "Tags". The "Tags" tab is active, showing a list of "Scanned Images". Each entry includes a version number, size (65 MB), and scan date. A red bar and a green bar indicate the scan status, with a warning icon and text stating "This image has vulnerabilities".

Image	Size	Scanned	Status
2	65 MB	Scanned 5 days ago	⚠ This image has vulnerabilities
2.4.1	65 MB	Scanned 9 days ago	⚠ This image has vulnerabilities
2.4	65 MB	Scanned 9 days ago	⚠ This image has vulnerabilities



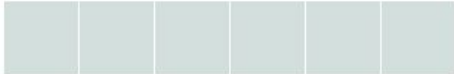
# Docker Security Scanning (Nautilus)

# Docker Security Scanning (Nautilus)

Scan results for **alpine:edge**



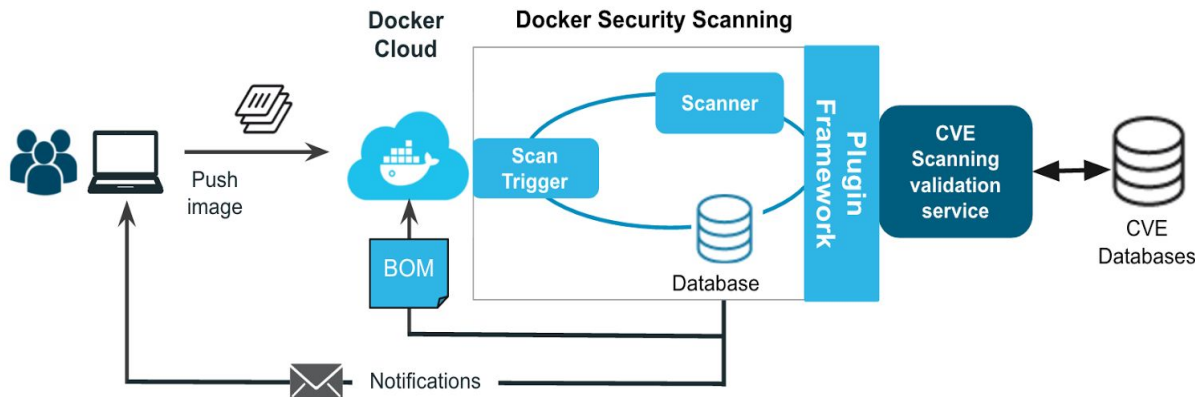
**Your image is clean!** No known vulnerabilities were found.  
Scanned 10 days ago

Layers	Components
1 <a href="#">ADD file:5cb42466c437...07c6a53832e650bc in /</a> 4.8MB  <div> No vulnerable components</div>	

<https://hub.docker.com/r/library/alpine/tags/>

- All official images on hub are scanned for vulnerabilities, lobby upstream for fixes!
- Can view scan results after logging into Docker Hub

# Docker Security Scanning (Nautilus)



- Checks against CVE database for declared layers
- Also performs binary scan to pick up on *statically linked binaries*
- Official repos have been scanned since Nov 2015, are rescanned often

# Hands-On Exercise

[github.com/riyazdf/dockercon-workshop](https://github.com/riyazdf/dockercon-workshop) - `trust` directory



# Capabilities

# Root vs Not Root

Capabilities breakdown root permissions into groups that can be individually allowed or blocked

- Often don't want or need all root permissions
- Can reduce attack surface by reducing capabilities

# Docker Default Capabilities

## In whitelist:

```
"CAP_CHOWN",  
"CAP_DAC_OVERRIDE",  
"CAP_FSETID",  
"CAP_FOWNER",  
"CAP_MKNOD",  
"CAP_NET_RAW",  
"CAP_SETGID",  
"CAP_SETUID",  
"CAP_SETFCAP",  
"CAP_SETPCAP",  
"CAP_NET_BIND_SERVICE",  
"CAP_SYS_CHROOT",  
"CAP_KILL",  
"CAP_AUDIT_WRITE",
```

## Not in whitelist:

```
"CAP_AUDIT_CONTROL", "CAP_AUDIT_READ",  
"CAP_BLOCK_SUSPEND", "CAP_DAC_READ_SEARCH",  
"CAP_IPC_LOCK", "CAP_IPC_OWNER",  
"CAP_LEASE", "CAP_LINUX_IMMUTABLE",  
"CAP_MAC_ADMIN", "CAP_MAC_OVERRIDE",  
"CAP_NET_ADMIN", "CAP_NET_BROADCAST",  
"CAP_SYS_ADMIN", "CAP_SYS_BOOT",  
"CAP_SYS_MODULE", "CAP_SYS_NICE",  
"CAP_SYS_PACCT", "CAP_SYS_PTRACE",  
"CAP_SYS_RAWIO", "CAP_SYS_RESOURCE",  
"CAP_SYS_TIME", "CAP_SYS_TTY_CONFIG",  
"CAP_SYSLOG", "CAP_WAKE_ALARM",
```

# How do we add/remove capabilities?

```
docker run --cap-add
```

```
docker run --cap-drop
```

```
docker run --cap-drop ALL --cap-add $CAP
```

# Configure capabilities in compose

cap\_add:

- CAP\_NET\_BROADCAST
- CAP\_NET\_RAW

cap\_drop:

- ALL

# What to watch out for

- Read the **fine print** for each capability!
  - man capabilities
  - i.e. removing CAP\_KILL only requires permissions checks and enabling bypasses permissions checks. It doesn't generically enable/disable the ability to kill
  - CAP\_SYS\_ADMIN is nearly root...

# What to watch out for

```
$ man capabilities
```

```
...
```

```
CAP_SYS_ADMIN
```

- \* Perform a range of system administration operations including: `quotactl(2)`, `mount(2)`, `umount(2)`, `swapon(2)`, `setdomainname(2)`;
- \* perform privileged `syslog(2)` operations (since Linux 2.6.37, `CAP_SYSLOG` should be used to permit such operations);
- \* perform `VM86_REQUEST_IRQ vm86(2)` command;
- \* perform `IPC_SET` and `IPC_RMID` operations on arbitrary System V IPC objects;
- \* override `RLIMIT_NPROC` resource limit;
- \* perform operations on *trusted* and *security* Extended Attributes (see `xattr(7)`);
- \* use `lookup_dcookie(2)`;
- \* use `ioprio_set(2)` to assign `IOPRIO_CLASS_RT` and (before Linux 2.6.25) `IOPRIO_CLASS_IDLE` I/O scheduling classes;
- \* forge PID when passing socket credentials via UNIX domain sockets;
- \* exceed `/proc/sys/fs/file-max`, the system-wide limit on the number of open files, in system calls that open files (e.g., `accept(2)`, `execve(2)`, `open(2)`, `pipe(2)`);
- \* employ `CLONE_*` flags that create new namespaces with `clone(2)` and `unshare(2)` (but, since Linux 3.8, creating user namespaces does not require any capability);
- \* call `perf_event_open(2)`;
- \* access privileged *perf* event information;
- \* call `perf_event_open(2)` for `CAP_SYS_ADMIN` in the kernel.

# Capabilities and Docker

- No extended attributes in images -> no capabilities elevation normally possible
- Use docker to reduce capabilities
- Docker can't grant capabilities to non-root users due to some limitations in older kernel versions



# Capabilities and Docker

Your options from worst to best:

1. **Manual management within the container:**  
`docker run --cap-add ALL`
2. **Restricted capabilities (still root):**  
`docker run --cap-drop ALL --cap-add ABC`
3. **No capabilities:**  
`docker run --user`

# What to watch out for



**\$ docker run --privileged ...**

gives **all capabilities** to the container, also lifts limitations from **device** cgroup

# Capabilities demo

# More information

[github.com/riyazdf/dockercon-workshop](https://github.com/riyazdf/dockercon-workshop) - `capabilities` directory

# Seccomp

# Original Seccomp

On-off feature that disabled all system calls except:

- `exit()`
- `read()`
- `write()`
- `sigreturn()`

# Seccomp-BPF

- Extension
- Allows us to configure what system calls are allowed/blocked
- Uses Berkeley Packet Filters (BPF)
- Allows examining system calls in detail before making a decision

# Is it enabled?

## In the kernel:

```
$ grep SECCOMP /boot/config-$(uname -r) # or zgrep SECCOMP /proc/config.gz  
CONFIG_SECCOMP=y  
CONFIG_SECCOMP_FILTER=y
```

## In docker:

```
$ docker run --rm alpine grep Seccomp /proc/self/status
```

## In docker 1.12:

```
$ docker info
```



# Default Whitelist

Lots of system calls, what's excluded:

acct	kexec_file_load	query_module	userfaultfd
add_key	kexec_load	quotactl	ustat
adjtimex	keyctl	reboot	vm86
bpf	lookup_dcookie	request_key	vm86old
clock_adjtime	mbind	set_mempolicy	
clock_settime	mount	setns	
clone	move_pages	settimeofday	
create_module	name_to_handle_at	stime	
delete_module	nfsservctl	swapon	
finit_module	open_by_handle_at	swapoff	
get_kernel_syms	perf_event_open	sysfs	
get_mempolicy	personality	_sysctl	
init_module	pivot_root	umount	
ioperm	process_vm_readv	umount2	
iopl	process_vm_writev	unshare	
kcmp	ptrace	uselib	

# The strace tool

```
$ strace -c -f -S name ls 2>&1 1>/dev/null | tail -n +3 | head -n -2 | awk '{print $(NF)}'
```

- access
- arch\_prctl
- brk
- close
- execve
- fstat
- getdents
- ioctl
- mmap
- mprotect
- munmap
- open
- read
- write

# Docker seccomp profile DSL

Seccomp policy example:

```
{
  "defaultAction": "SCMP_ACT_ERRNO",
  "architectures": [
    "SCMP_ARCH_X86_64",
    "SCMP_ARCH_X86",
    "SCMP_ARCH_X32"
  ],
  "syscalls": [
    {
      "name": "accept",
      "action": "SCMP_ACT_ALLOW",
      "args": []
    },
    ...
  ]
}
```

Possible actions:

```
SCMP_ACT_KILL
SCMP_ACT_TRAP
SCMP_ACT_ERRNO
SCMP_ACT_TRACE
SCMP_ACT_ALLOW
```

# Docker seccomp profile DSL

More complex filters:

```
"args": [  
  {  
    "index": 0,  
    "op": "SCMP_CMP_MASKED_EQ",  
    "value": 2080505856,  
    "valueTwo": 0  
  }  
]
```

*2080505856 == 0x7C020000*

# Seccomp and the no-new-privileges option



Seccomp policies have to be applied before executing your container and be less specific unless you use:

```
--security-opt no-new-privileges
```

In this case you need to allow only `futex` `stat` `execve`.

This flag also disables `setuid` binaries:

```
$ sudo ls
```

```
sudo: effective uid is not 0, is /usr/bin/sudo on a file system with the 'nosuid' option set  
or an NFS file system without root privileges?
```

# More information

[github.com/riyazdf/dockercon-workshop](https://github.com/riyazdf/dockercon-workshop) - `seccomp` directory

```
$ docker run --rm -it --security-opt seccomp=default-no-chmod.json alpine  
chmod 777 /  
chmod: /: Operation not permitted
```

# Linux Security Modules

# What is a LSM?

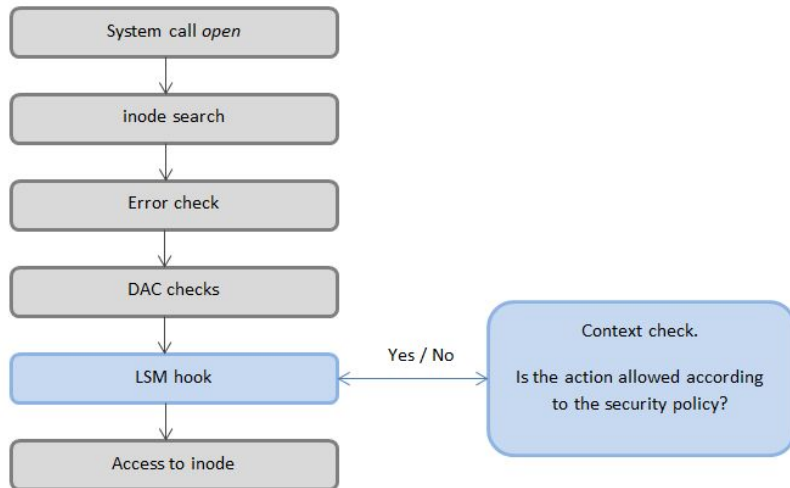
A plugin to the linux kernel that allows us to set policies to restrict what a process can do.

**Mandatory Access Control:** instead of using user-defined permissions to specify access, the underlying system describes permissions itself with labels



# What is a LSM?

**Under the hood:** each LSM implements a kernel interface that hooks into user-level syscalls about to access an important kernel object (inodes, task control blocks, etc.), either allowing them to pass through or denying them outright depending on the application profile



# Available LSMs

AppArmor



SELinux



Smack

Tomoyo

# Deep Dive - AppArmor

## File Access Management

AppArmor uses globbing and deny syntax to express filepath restrictions

- `deny /sys/* rwklx` Deny read/write/lock/link/execute on files in /sys/
- `deny /sys/** rwklx` Deny on files in /sys/ **and subdirectories**

# Deep Dive - AppArmor

## Networking Management

*Like firewall rules:*

- Can completely disable networking:
- Can deny certain permissions:
- Can specify specific IP/ports :

```
deny network
```

```
deny network bind, inet
```

```
network tcp src 192.168.1.1:80 dst 170.1.1.0:80
```

# Deep Dive - AppArmor

## Capability Management

AppArmor can also deny capabilities with a simple syntax:

- `deny capability chown,`
- `deny capability dac_override`

# Deep Dive - AppArmor

## Composability

### C-style include statements

- `include <abstractions/base>` - built-in bundle of files
- `include "/etc/apparmor.d/include/foo"` - absolute path from file
- `include "../relative_path/bar"` - relative path from file

# Deep Dive - AppArmor

Tools for debugging and generating profiles (on Ubuntu):

```
$ sudo apt install apparmor-utils
```

```
$ aa-complain <PATH_TO_PROFILE> # Watch AppArmor block things!
```

```
$ aa-genprof <PATH_TO_BINARY> # Interactive profile generation!
```

# Do I still need Seccomp and Cap-drop?

**Why not?** Docker sets a profile for each setting by default

- Some overlap but each feature still adds unique functionality
- Defense-in-depth



# Common mistake: disabling profiles



**SELinux:** `setenforce 0` (on daemon)

<http://stopdisablinglinux.com/>

**AppArmor:** `--security-opt apparmor:unconfined` (on docker run)

*This one's a little harder to do "by accident"*

`docker run --privileged`

# Hands-On Exercise

[github.com/riyazdf/dockercon-workshop](https://github.com/riyazdf/dockercon-workshop) - `apparmor` directory

# Docker Bench

<https://dockerbench.com>

- Open-source tool for running automated tests
  - inspired by the CIS Docker 1.11 benchmark
- Runs against containers currently running on same host
- Checks for AppArmor, read-only volumes, etc...

```
# -----
# Docker Bench for Security v1.0.0
#
# Docker, Inc. (c) 2015-
#
# Checks for dozens of common best-practices around deploying Docker containers in production.
# Inspired by the CIS Docker 1.11 Benchmark:
# https://benchmarks.cisecurity.org/downloads/show-single/index.cfm?file=docker16.110
# -----

Initializing Sat Apr 30 23:04:50 CEST 2016

[INFO] 1 - Host Configuration
[WARN] 1.1 - Create a separate partition for containers
[PASS] 1.2 - Use an updated Linux Kernel
[PASS] 1.4 - Remove all non-essential services from the host - Network
[PASS] 1.5 - Keep Docker up to date
[INFO] * Using 1.12.0 which is current as of 2016-04-27
[INFO] * Check with your operating system vendor for support and security maintenance for docker
[INFO] 1.6 - Only allow trusted users to control Docker daemon
[INFO] * docker:x:999:tsj
[PASS] 1.7 - Audit docker daemon - /usr/bin/docker
[PASS] 1.8 - Audit Docker files and directories - /var/lib/docker
[PASS] 1.9 - Audit Docker files and directories - /etc/docker
[PASS] 1.10 - Audit Docker files and directories - docker.service
[PASS] 1.11 - Audit Docker files and directories - docker.socket
[PASS] 1.12 - Audit Docker files and directories - /etc/default/docker
[INFO] 1.13 - Audit Docker files and directories - /etc/docker/daemon.json
[INFO] * File not found
[PASS] 1.14 - Audit Docker files and directories - /usr/bin/docker-containerd
[PASS] 1.15 - Audit Docker files and directories - /usr/bin/docker-runc

[INFO] 2 - Docker Daemon Configuration
[PASS] 2.1 - Restrict network traffic between containers
[PASS] 2.2 - Set the logging level
[PASS] 2.3 - Allow Docker to make changes to iptables
[PASS] 2.4 - Do not use insecure registries
[PASS] 2.5 - Do not use the aufs storage driver
[INFO] 2.6 - Configure TLS authentication for Docker daemon
[INFO] * Docker daemon not listening on TCP
[INFO] 2.7 - Set default ulimit as appropriate
[INFO] * Default ulimit doesn't appear to be set
[WARN] 2.8 - Enable user namespace support
[PASS] 2.9 - Confirm default cgroup usage
[PASS] 2.10 - Do not change base device size until needed
[WARN] 2.11 - Use authorization plugin
[WARN] 2.12 - Configure centralized and remote logging
[PASS] 2.13 - Disable operations on legacy registry (v1)
```

# *View from 10,000 feet:*

## **Docker Security Checklist**

### **Build:**

- Use minimal images (alpine)
- Use official images
- Using images pulled by content trust (fresh, pulled by digest from authors you trust)

### **Ship:**

- Push to your consumers with content trust
- View results from Docker Security Scanning

### **Run:**

- Mutual TLS between client/engine
- Read-only volumes and containers
- User namespaces in the daemon
- Limit resources with cgroups
- Use the default apparmor/seccomp/capabilities, or your own tested profiles (not --privileged!)

# Thank you!

- Please take the exercises home! We'll accept issues and pull requests :)
- Learn more at <https://docs.docker.com/engine/security/security/>

# Advanced Topics

“Extra for Experts”

# AuthZ plugins

# Running your own Notary

Deploy a notary

```
$ git clone https://github.com/docker/notary.git
```

```
$ cd notary
```

```
$ docker-compose up
```



# Notary Delegations

```
(admin)$ notary key rotate <GUN> snapshot -r  
(user)$ < generates private key and x509 cert, gives user.crt to admin >  
(admin)$ notary delegation add <GUN> targets/user user.crt --all-paths  
(admin)$ notary publish <GUN>
```

Docker engine  $\geq$  1.11 will sign with delegation keys if it detects them

